

Pandas Tutorial

I. Overview

This tutorial walks through the fundamentals of the Pandas data science library.

II. Preparation

For this tutorial, you will need the Pandas library for data manipulation as well as the Seaborn library for visualization. Then you will need to download the data.

INSTALL THE LIBRARIES

There are two primary ways to install these libraries. The first is to use a package manager that is included in an IDE such as Thonny. The second is to use a command line install tool, like PIP.

If you are using an IDE with a package manager, such as Thonny, select Tools -> Manage Packages. Search for **pandas** and install it. Then, search for **seaborn** and install it as well.

If you are installing these libraries from the command line, you can type **pip install pandas** and **pip install seaborn** . To install them.

Instructor Tip

If you are trying to use the pip installer, you will likely need to have administrative privileges by running "sudo pip install pandas" on Mac/Linux or, if you are using Windows, you can right click the Command Line utility and select "Run as Administrator" rather than running it directly.

If you are trying to use the pip installer and encounter an error message similar to a "command not found" it is likely that the directory containing this utility is not in your path (the place where your computer looks for applications to run). You may need to find the pip utility yourself and then add the folder that contains it, to the path.

DOWNLOADING THE DATASET

This tutorial uses a dataset of NBA basketball statistics that can be downloaded [here](#). This dataset was originally obtained from opensourcesports.com and can also be found at [Kaggle](#). It did not come with an explicit license, but based on other datasets from Open Source Sports, we treat it as follows: This database is copyright 1996-2015 by Sean Lahman. This work is licensed under a Creative Commons Attribution-ShareAlike 3.0 Unported License. For details see: <http://creativecommons.org/licenses/by-sa/3.0/>

The dataset is contained in several files contained in a ZIP file. You should download and extract the ZIP file, then you can copy the .csv (comma separated values) files contained in the ZIP file into another directory to work with. It is most convenient to put these .csv files in the same directory as your .py python script.

The two primary files we will use from this dataset are:

- » basketball_players.csv - This file contains the stats for each player for a given season.
- » basketball_master.csv - This file contains additional information about the players, such as biographical information, etc.

III. Loading the libraries and data files in Python

The first step is to load the libraries.

```
import pandas as pd # Our data manipulation library
import seaborn as sns # Used for graphing/plotting
import matplotlib.pyplot as plt # If we need any low level methods
import os # Used to change the directory to the right place
```

The next step is to load the data itself.

```
# This line isn't necessary, but it makes it so the later commands (e.g., read_csv)
# are in a consistent place (You will obviously need to change this to the correct location on _
# If you have put the data files and your Python script in the same folder, you
# don't need this line.
os.chdir("/Users/sburton/cs241/nbaData/")

# Load in the data
# The players data (basketball_players.csv) has the season stats
players = pd.read_csv("basketball_players.csv")
```

After running the above commands, I received the following message:

```
/Library/Frameworks/Python.framework/Versions/3.6/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2717: DtypeWarning: Columns (41) have mixed types. Specify dtype option on import or set low_memory=False.
interactivity=interactivity, compiler=compiler, result=result)
```

The above error message says that some column data types could not be determined automatically, or seemed to be mixed. We can avoid this by specifying them directly, but in this case, we will ignore it, and move on to look at the data.

With the players dataset loaded, we can play with it to see different things.

```
players
```

Instructor Tip:

If you are running the commands directly in the interactive Python console, you can just type an expression or the name of the dataset variable directly and it will display the results. If you are putting it into a script, you'll need to print the results (e.g., "print(players)") rather than just typing the name of the variable.

| | playerID | year | stint | tmID | lgID | GP | GS | minutes | points | oRebounds | \ |
|-------|------------|------|---------------|------|--------|-----|-----------------|---------|------------|-----------|---|
| 0 | abramjo01 | 1946 | 1 | PIT | NBA | 47 | 0 | 0 | 527 | 0 | |
| 1 | aubucch01 | 1946 | 1 | DTF | NBA | 30 | 0 | 0 | 65 | 0 | |
| 2 | bakerno01 | 1946 | 1 | CHS | NBA | 4 | 0 | 0 | 0 | 0 | |
| 3 | baltiie01 | 1946 | 1 | STB | NBA | 58 | 0 | 0 | 138 | 0 | |
| 4 | barrjo01 | 1946 | 1 | STB | NBA | 58 | 0 | 0 | 295 | 0 | |
| 5 | baumhfr01 | 1946 | 1 | CLR | NBA | 45 | 0 | 0 | 631 | 0 | |
| 6 | beckemo01 | 1946 | 1 | PIT | NBA | 17 | 0 | 0 | 108 | 0 | |
| 7 | beckemo01 | 1946 | 2 | BOS | NBA | 6 | 0 | 0 | 13 | 0 | |
| 8 | beckemo01 | 1946 | 3 | DTF | NBA | 20 | 0 | 0 | 41 | 0 | |
| 9 | beendha01 | 1946 | 1 | PRO | NBA | 58 | 0 | 0 | 713 | 0 | |
| 10 | biasaha01 | 1946 | 1 | TRH | NBA | 6 | 0 | 0 | 6 | 0 | |
| 11 | brighal01 | 1946 | 1 | BOS | NBA | 58 | 0 | 0 | 567 | 0 | |
| 12 | brindau01 | 1946 | 1 | NYK | NBA | 12 | 0 | 0 | 34 | 0 | |
| 13 | brownha01 | 1946 | 1 | DTF | NBA | 54 | 0 | 0 | 264 | 0 | |
| 14 | brownle01 | 1946 | 1 | CLR | NBA | 5 | 0 | 0 | 0 | 0 | |
| 15 | byrneto01 | 1946 | 1 | NYK | NBA | 60 | 0 | 0 | 453 | 0 | |
| 16 | bytzum01 | 1946 | 1 | PIT | NBA | 60 | 0 | 0 | 210 | 0 | |
| 17 | callato01 | 1946 | 1 | PRO | NBA | 13 | 0 | 0 | 17 | 0 | |
| 18 | calveer01 | 1946 | 1 | PRO | NBA | 59 | 0 | 0 | 845 | 0 | |
| 19 | carlich01 | 1946 | 1 | CHS | NBA | 51 | 0 | 0 | 256 | 0 | |
| 20 | carlsdo01 | 1946 | 1 | CHS | NBA | 59 | 0 | 0 | 630 | 0 | |
| 21 | cluggbo01 | 1946 | 1 | NYK | NBA | 54 | 0 | 0 | 238 | 0 | |
| 22 | connoch01 | 1946 | 1 | BOS | NBA | 49 | 0 | 0 | 227 | 0 | |
| 23 | corleke01 | 1946 | 1 | CLR | NBA | 3 | 0 | 0 | 0 | 0 | |
| 24 | crisilha01 | 1946 | 1 | BOS | NBA | 4 | 0 | 0 | 6 | 0 | |
| 25 | curear01 | 1946 | 1 | PRO | NBA | 12 | 0 | 0 | 10 | 0 | |
| 26 | dallmho01 | 1946 | 1 | PHW | NBA | 60 | 0 | 0 | 528 | 0 | |
| 27 | davisau01 | 1946 | 1 | STB | NBA | 59 | 0 | 0 | 287 | 0 | |
| 28 | davisbi01 | 1946 | 1 | CHS | NBA | 47 | 0 | 0 | 84 | 0 | |
| 29 | dehnehe01 | 1946 | 1 | PRO | NBA | 10 | 0 | 0 | 14 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 23721 | lacoufr01 | 1962 | 0 | OA0 | ABL1 | 25 | 0 | 901 | 494 | 0 | |
| 23722 | millsda01 | 1962 | 0 | OA0 | ABL1 | 21 | 0 | 359 | 155 | 0 | |
| 23723 | romanni01 | 1962 | 0 | OA0 | ABL1 | 25 | 0 | 755 | 225 | 0 | |
| 23724 | toddro01 | 1962 | 0 | OA0 | ABL1 | 24 | 0 | 733 | 339 | 0 | |
| 23725 | turneja02 | 1962 | 0 | OA0 | ABL1 | 25 | 0 | 930 | 373 | 0 | |
| 23726 | vaughgo01 | 1962 | 0 | OA0 | ABL1 | 22 | 0 | 350 | 122 | 0 | |
| 23727 | wilkibo01 | 1962 | 0 | OA0 | ABL1 | 23 | 0 | 339 | 128 | 0 | |
| 23728 | yateswa01 | 1962 | 0 | OA0 | ABL1 | 25 | 0 | 624 | 268 | 0 | |
| 23729 | bellwh01 | 1962 | 0 | PGR | ABL1 | 7 | 0 | 223 | 69 | 0 | |
| 23730 | curtich01 | 1962 | 0 | PGR | ABL1 | 22 | 0 | 666 | 301 | 0 | |
| 23731 | curtihe01 | 1962 | 0 | PGR | ABL1 | 20 | 0 | 596 | 275 | 0 | |
| 23732 | hawkico01 | 1962 | 0 | PGR | ABL1 | 16 | 0 | 668 | 447 | 0 | |
| 23733 | kennewa01 | 1962 | 0 | PGR | ABL1 | 5 | 0 | 91 | 50 | 0 | |
| 23734 | manghwa01 | 1962 | 0 | PGR | ABL1 | 22 | 0 | 529 | 213 | 0 | |
| 23735 | mccarjo01 | 1962 | 0 | PGR | ABL1 | 18 | 0 | 420 | 95 | 0 | |
| 23736 | mccoyji01 | 1962 | 0 | PGR | ABL1 | 22 | 0 | 847 | 350 | 0 | |
| 23737 | rolliph01 | 1962 | 0 | PGR | ABL1 | 17 | 0 | 611 | 221 | 0 | |
| 23738 | tyrach01 | 1962 | 0 | PGR | ABL1 | 23 | 0 | 518 | 212 | 0 | |
| 23739 | wiesebo01 | 1962 | 0 | PGR | ABL1 | 15 | 0 | 252 | 101 | 0 | |
| 23740 | alixlo01 | 1962 | 0 | PHT | ABL1 | 1 | 0 | 3 | 0 | 0 | |
| 23741 | blyesy01 | 1962 | 0 | PHT | ABL1 | 28 | 0 | 975 | 496 | 0 | |
| 23742 | chmiemo01 | 1962 | 0 | PHT | ABL1 | 20 | 0 | 528 | 208 | 0 | |
| 23743 | clarkbo01 | 1962 | 0 | PHT | ABL1 | 16 | 0 | 217 | 59 | 0 | |
| 23744 | hillcl01 | 1962 | 0 | PHT | ABL1 | 22 | 0 | 422 | 145 | 0 | |
| 23745 | johnsan01 | 1962 | 0 | PHT | ABL1 | 28 | 0 | 732 | 314 | 0 | |
| 23746 | kaisero01 | 1962 | 0 | PHT | ABL1 | 27 | 0 | 978 | 467 | 0 | |
| 23747 | spragbr01 | 1962 | 0 | PHT | ABL1 | 27 | 0 | 746 | 356 | 0 | |
| 23748 | tayloro02 | 1962 | 0 | PHT | ABL1 | 28 | 0 | 1007 | 355 | 0 | |
| 23749 | wellsra01 | 1962 | 0 | PHT | ABL1 | 2 | 0 | 36 | 4 | 0 | |
| 23750 | wrighle01 | 1962 | 0 | PHT | ABL1 | 28 | 0 | 813 | 195 | 0 | |
| ... | PostBlocks | | PostTurnovers | | PostPF | | PostfgAttempted | | PostfgMade | \ | |
| 0 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 1 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 2 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 3 | ... | 0 | | 0 | 3 | | 10 | | 2 | | |
| 4 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 5 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 6 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 7 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 8 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 9 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 10 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 11 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 12 | ... | 0 | | 0 | 4 | | 6 | | 3 | | |
| 13 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 14 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 15 | ... | 0 | | 0 | 2 | | 46 | | 11 | | |
| 16 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 17 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 18 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 19 | ... | 0 | | 0 | 33 | | 88 | | 20 | | |
| 20 | ... | 0 | | 0 | 31 | | 200 | | 54 | | |
| 21 | ... | 0 | | 0 | 12 | | 27 | | 4 | | |
| 22 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 23 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 24 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 25 | ... | 0 | | 0 | 0 | | 0 | | 0 | | |
| 26 | ... | 0 | | 0 | 28 | | 104 | | 26 | | |

```
27    ...      0      0      3      6      2
28    ...      0      0     10     14      2
29    ...      0      0      0      0      0
...
23721 ...      0      0      0      0      0
23722 ...      0      0      0      0      0
23723 ...      0      0      0      0      0
23724 ...      0      0      0      0      0
23725 ...      0      0      0      0      0
23726 ...      0      0      0      0      0
23727 ...      0      0      0      0      0
23728 ...      0      0      0      0      0
23729 ...      0      0      0      0      0
23730 ...      0      0      0      0      0
23731 ...      0      0      0      0      0
23732 ...      0      0      0      0      0
23733 ...      0      0      0      0      0
23734 ...      0      0      0      0      0
23735 ...      0      0      0      0      0
23736 ...      0      0      0      0      0
23737 ...      0      0      0      0      0
23738 ...      0      0      0      0      0
23739 ...      0      0      0      0      0
23740 ...      0      0      0      0      0
23741 ...      0      0      0      0      0
23742 ...      0      0      0      0      0
23743 ...      0      0      0      0      0
23744 ...      0      0      0      0      0
23745 ...      0      0      0      0      0
23746 ...      0      0      0      0      0
23747 ...      0      0      0      0      0
23748 ...      0      0      0      0      0
23749 ...      0      0      0      0      0
23750 ...      0      0      0      0      0
      PostftAttempted  PostftMade  PostthreeAttempted  PostthreeMade  note
0                    0            0                    0              0    NaN
1                    0            0                    0              0    NaN
2                    0            0                    0              0    NaN
3                    1            0                    0              0    NaN
4                    0            0                    0              0    NaN
5                    0            0                    0              0    NaN
6                    0            0                    0              0    NaN
7                    0            0                    0              0    NaN
8                    0            0                    0              0    NaN
9                    0            0                    0              0    NaN
10                   0            0                    0              0    NaN
11                   0            0                    0              0    NaN
12                   6            4                    0              0    NaN
13                   0            0                    0              0    NaN
14                   0            0                    0              0    NaN
15                  11            2                    0              0    NaN
16                   0            0                    0              0    NaN
17                   0            0                    0              0    NaN
18                   0            0                    0              0    NaN
19                  28           16                    0              0    NaN
20                  44           27                    0              0    NaN
21                   2            0                    0              0    NaN
22                   0            0                    0              0    NaN
23                   0            0                    0              0    NaN
24                   0            0                    0              0    NaN
25                   0            0                    0              0    NaN
26                  40           30                    0              0    NaN
27                   3            3                    0              0    NaN
28                   5            2                    0              0    NaN
29                   0            0                    0              0    NaN
...
23721 ...      0      0      0      0      NaN
23722 ...      0      0      0      0      NaN
23723 ...      0      0      0      0      NaN
23724 ...      0      0      0      0      NaN
23725 ...      0      0      0      0      NaN
23726 ...      0      0      0      0      NaN
23727 ...      0      0      0      0      NaN
23728 ...      0      0      0      0      NaN
23729 ...      0      0      0      0      NaN
23730 ...      0      0      0      0      NaN
23731 ...      0      0      0      0      NaN
23732 ...      0      0      0      0      NaN
23733 ...      0      0      0      0      NaN
23734 ...      0      0      0      0      NaN
23735 ...      0      0      0      0      NaN
23736 ...      0      0      0      0      NaN
23737 ...      0      0      0      0      NaN
23738 ...      0      0      0      0      NaN
23739 ...      0      0      0      0      NaN
23740 ...      0      0      0      0      NaN
23741 ...      0      0      0      0      NaN
23742 ...      0      0      0      0      NaN
23743 ...      0      0      0      0      NaN
23744 ...      0      0      0      0      NaN
23745 ...      0      0      0      0      NaN
23746 ...      0      0      0      0      NaN
23747 ...      0      0      0      0      NaN
23748 ...      0      0      0      0      NaN
23749 ...      0      0      0      0      NaN
23750 ...      0      0      0      0      NaN

[23751 rows x 42 columns]
```

IV. Exploring the Data

First we might explore this data a little bit to see what we have. We can see the available columns:

```
players.columns

Index(['playerID', 'year', 'stint', 'tmID', 'lgID', 'GP', 'GS', 'minutes',
      'points', 'oRebounds', 'dRebounds', 'rebounds', 'assists', 'steals',
      'blocks', 'turnovers', 'PF', 'fgAttempted', 'fgMade', 'ftAttempted',
      'ftMade', 'threeAttempted', 'threeMade', 'PostGP', 'PostGS',
      'PostMinutes', 'PostPoints', 'PostoRebounds', 'PostdRebounds',
      'PostRebounds', 'PostAssists', 'PostSteals', 'PostBlocks',
      'PostTurnovers', 'PostPF', 'PostfgAttempted', 'PostfgMade',
      'PostftAttempted', 'PostftMade', 'PostthreeAttempted', 'PostthreeMade',
```

```
'note'],
dtype='object')
```

And we can look at statistics about certain variables. For example, we can look at the min, max, mean, and median for a column like rebounds:

```
min = players["rebounds"].min()
max = players["rebounds"].max()
mean = players["rebounds"].mean()
median = players["rebounds"].median()

print("Rebounds per season: Min:{}, Max:{}, Mean:{:.2f}, Median:{}".format(min, max, mean, median))
```

Rebounds per season: Min:0, Max:2149, Mean:209.06, Median:133.0

Instructor Tip:

When working with existing columns, you can either use the dot notation "players.rebounds" or the square bracket notation "players["rebounds"]". If you are creating a new column or if your column name has a space in it, you must use the square bracket notation.

FINDING THE BEST REBOUNDERS

Perhaps we want to look at the highest rebounding seasons to see the player that had that amount and the team they played on. We can sort the data by rebounds and print out the top 10 rows:

```
print(players.sort_values("rebounds", ascending=False).head(10))
```

| | playerID | year | stint | tmID | lgID | GP | GS | minutes | points | oRebounds | \ |
|------------------------|-----------------|---------------|--------------------|-----------------|------------|-----|----|---------|--------|-----------|---|
| 1972 | chambwi01 | 1960 | 1 | PHW | NBA | 79 | 0 | 3773 | 3033 | 0 | |
| 2078 | chambwi01 | 1961 | 1 | PHW | NBA | 80 | 0 | 3882 | 4029 | 0 | |
| 2697 | chambwi01 | 1966 | 1 | PHI | NBA | 81 | 0 | 3682 | 1956 | 0 | |
| 2856 | chambwi01 | 1967 | 1 | PHI | NBA | 82 | 0 | 3836 | 1992 | 0 | |
| 2199 | chambwi01 | 1962 | 1 | SFW | NBA | 80 | 0 | 3806 | 3586 | 0 | |
| 2578 | chambwi01 | 1965 | 1 | PHI | NBA | 79 | 0 | 3737 | 2649 | 0 | |
| 1859 | chambwi01 | 1959 | 1 | PHW | NBA | 72 | 0 | 3338 | 2707 | 0 | |
| 2403 | russebi01 | 1963 | 1 | BOS | NBA | 78 | 0 | 3482 | 1168 | 0 | |
| 2534 | russebi01 | 1964 | 1 | BOS | NBA | 78 | 0 | 3466 | 1102 | 0 | |
| 2043 | russebi01 | 1960 | 1 | BOS | NBA | 78 | 0 | 3458 | 1322 | 0 | |
| ... | PostBlocks | PostTurnovers | PostPF | PostfgAttempted | PostfgMade | \ | | | | | |
| 1972 | ... | 0 | 0 | 10 | 96 | 45 | | | | | |
| 2078 | ... | 0 | 0 | 27 | 347 | 162 | | | | | |
| 2697 | ... | 0 | 0 | 37 | 228 | 132 | | | | | |
| 2856 | ... | 0 | 0 | 29 | 232 | 124 | | | | | |
| 2199 | ... | 0 | 0 | 0 | 0 | 0 | | | | | |
| 2578 | ... | 0 | 0 | 10 | 110 | 56 | | | | | |
| 1859 | ... | 0 | 0 | 17 | 252 | 125 | | | | | |
| 2403 | ... | 0 | 0 | 23 | 132 | 47 | | | | | |
| 2534 | ... | 0 | 0 | 43 | 150 | 79 | | | | | |
| 2043 | ... | 0 | 0 | 24 | 171 | 73 | | | | | |
| | PostftAttempted | PostftMade | PostthreeAttempted | PostthreeMade | note | | | | | | |
| 1972 | 38 | 21 | 0 | 0 | NaN | | | | | | |
| 2078 | 151 | 96 | 0 | 0 | NaN | | | | | | |
| 2697 | 160 | 62 | 0 | 0 | NaN | | | | | | |
| 2856 | 158 | 60 | 0 | 0 | NaN | | | | | | |
| 2199 | 0 | 0 | 0 | 0 | NaN | | | | | | |
| 2578 | 68 | 28 | 0 | 0 | NaN | | | | | | |
| 1859 | 110 | 49 | 0 | 0 | NaN | | | | | | |
| 2403 | 67 | 37 | 0 | 0 | NaN | | | | | | |
| 2534 | 76 | 40 | 0 | 0 | NaN | | | | | | |
| 2043 | 86 | 45 | 0 | 0 | NaN | | | | | | |
| [10 rows x 42 columns] | | | | | | | | | | | |

That is showing a lot of columns and making it hard to read, so we might repeat it and only show a few:

```
print(players[["playerID", "year", "tmID", "rebounds"]].sort_values("rebounds", ascending=False))
```

| | playerID | year | tmID | rebounds |
|------|-----------|------|------|----------|
| 1972 | chambwi01 | 1960 | PHW | 2149 |
| 2078 | chambwi01 | 1961 | PHW | 2052 |
| 2697 | chambwi01 | 1966 | PHI | 1957 |
| 2856 | chambwi01 | 1967 | PHI | 1952 |
| 2199 | chambwi01 | 1962 | SFW | 1946 |
| 2578 | chambwi01 | 1965 | PHI | 1943 |
| 1859 | chambwi01 | 1959 | PHW | 1941 |
| 2403 | russebi01 | 1963 | BOS | 1930 |
| 2534 | russebi01 | 1964 | BOS | 1878 |
| 2043 | russebi01 | 1960 | BOS | 1868 |

MERGING OR JOINING SEPARATE DATASETS

The previous results are certainly much easier to view, however, while the player ID may help us know the name of the player, to get their actual name and other biographical information, we need to load in another dataset (the "master" dataset) and connect them. When we connect them, we do a Left Join which says that we want every row in our players dataset, even if the master dataset doesn't have information for them. Then we need to specify that the columns that match them up

(playerID from the players dataset matches bioID from the master dataset

```
# The "master" data (basketball_master.csv) has names, biographical information, etc.
master = pd.read_csv("basketball_master.csv")

# We can do a left join to "merge" these two datasets together
nba = pd.merge(players, master, how="left", left_on="playerID", right_on="bioID")
```

At this point the variable `nba` contains a full dataset with many different rows and columns. By printing that variable or its columns, you can see a summary of the dataset.

```
print(nba.columns)

Index(['playerID', 'year', 'stint', 'tmID', 'lgID', 'GP', 'GS', 'minutes',
      'points', 'oRebounds', 'dRebounds', 'rebounds', 'assists', 'steals',
      'blocks', 'turnovers', 'PF', 'fgAttempted', 'fgMade', 'ftAttempted',
      'ftMade', 'threeAttempted', 'threeMade', 'PostGP', 'PostGS',
      'PostMinutes', 'PostPoints', 'PostoRebounds', 'PostdRebounds',
      'PostRebounds', 'PostAssists', 'PostSteals', 'PostBlocks',
      'PostTurnovers', 'PostPF', 'PostfgAttempted', 'PostfgMade',
      'PostftAttempted', 'PostftMade', 'PostthreeAttempted', 'PostthreeMade',
      'note', 'bioID', 'useFirst', 'firstName', 'middleName', 'lastName',
      'nameGiven', 'fullGivenName', 'nameSuffix', 'nameNick', 'pos',
      'firstseason', 'lastseason', 'height', 'weight', 'college',
      'collegeOther', 'birthDate', 'birthCity', 'birthState', 'birthCountry',
      'highSchool', 'hsCity', 'hsState', 'hsCountry', 'deathDate', 'race'],
      dtype='object')
```

Notice that we have all the columns from the players dataset before, but we have also added a number of columns such as first and last names, college, birth city, etc. It turns out that many people don't use their given first name as the name they go by, so this dataset also has a `useFirst` column that is the first name they use.

With this additional information, we can return to printing out the top rebounders, but this time, we will use the `nba` variable and use different columns:

```
print(nba[["year", "useFirst", "lastName", "tmID", "rebounds"]].sort_values("rebounds", ascending=False))
```

| | year | useFirst | lastName | tmID | rebounds |
|------|------|----------|-------------|------|----------|
| 1972 | 1960 | Wilt | Chamberlain | PHW | 2149 |
| 2078 | 1961 | Wilt | Chamberlain | PHW | 2052 |
| 2697 | 1966 | Wilt | Chamberlain | PHI | 1957 |
| 2856 | 1967 | Wilt | Chamberlain | PHI | 1952 |
| 2199 | 1962 | Wilt | Chamberlain | SFW | 1946 |
| 2578 | 1965 | Wilt | Chamberlain | PHI | 1943 |
| 1859 | 1959 | Wilt | Chamberlain | PHW | 1941 |
| 2403 | 1963 | Bill | Russell | BOS | 1930 |
| 2534 | 1964 | Bill | Russell | BOS | 1878 |
| 2043 | 1960 | Bill | Russell | BOS | 1868 |

This is much easier to see the players in this result. We could go through a similar process to better understand the `tmID` (team ID) by joining in another dataset, but that will not be explored in this tutorial.

CREATING NEW COLUMNS

While the total number of rebounds in a season is interesting, most people like to compare an average of rebounds per game. Unfortunately, a column for this is not available in this dataset, however, there is a column for rebounds, and one for games played (`GP`) so we can make a new column on our own that contains this information by dividing rebounds by games played:

```
nba["reboundsPerGame"] = nba["rebounds"] / nba["GP"]
print(nba[["year", "useFirst", "lastName", "rebounds", "GP", "reboundsPerGame"]].sort_values("reboundsPerGame", ascending=False))
```

| | year | useFirst | lastName | rebounds | GP | reboundsPerGame |
|-------|------|----------|-------------|----------|----|-----------------|
| 19295 | 2007 | Mike | Conley | 118 | 0 | inf |
| 19828 | 2008 | Mike | Conley | 251 | 0 | inf |
| 1972 | 1960 | Wilt | Chamberlain | 2149 | 79 | 27.202532 |
| 1859 | 1959 | Wilt | Chamberlain | 1941 | 72 | 26.958333 |
| 2078 | 1961 | Wilt | Chamberlain | 2052 | 80 | 25.650000 |
| 2403 | 1963 | Bill | Russell | 1930 | 78 | 24.743590 |
| 2578 | 1965 | Wilt | Chamberlain | 1943 | 79 | 24.594937 |
| 2199 | 1962 | Wilt | Chamberlain | 1946 | 80 | 24.325000 |
| 2697 | 1966 | Wilt | Chamberlain | 1957 | 81 | 24.160494 |
| 2534 | 1964 | Bill | Russell | 1878 | 78 | 24.076923 |

This gives us a nice view of the top performers for rebounds per game. As we look at the dataset, we should see if this makes sense. The first thing that should stand out is that Mike Conley is listed as having

rebounds in 2007 and 2008 without playing any games. This must clearly be an error in the data. It turns out this kind of thing happens all the time in real datasets, so we always need to be on the lookout for it. If we have the ability, we could talk to the people that produced the data and see if they could fix the problem, but in this case, we don't have that option. So our best choice is probably to remove any rows that don't have games played.

```
# Let's just remove any rows with GP=0
nba = nba[nba.GP > 0]
```

Now, try the last commands again:

```
nba["reboundsPerGame"] = nba["rebounds"] / nba["GP"]
print(nba[["year", "useFirst", "lastName", "rebounds", "GP", "reboundsPerGame"]].sort_values("reboundsPerGame"))
```


| | year | useFirst | lastName | rebounds | GP | reboundsPerGame |
|------|------|----------|-------------|----------|----|-----------------|
| 1972 | 1960 | Wilt | Chamberlain | 2149 | 79 | 27.202532 |
| 1859 | 1959 | Wilt | Chamberlain | 1941 | 72 | 26.958333 |
| 2078 | 1961 | Wilt | Chamberlain | 2052 | 80 | 25.650000 |
| 2403 | 1963 | Bill | Russell | 1930 | 78 | 24.743590 |
| 2578 | 1965 | Wilt | Chamberlain | 1943 | 79 | 24.594937 |
| 2199 | 1962 | Wilt | Chamberlain | 1946 | 80 | 24.325000 |
| 2697 | 1966 | Wilt | Chamberlain | 1957 | 81 | 24.160494 |
| 2534 | 1964 | Bill | Russell | 1878 | 78 | 24.076923 |
| 1936 | 1959 | Bill | Russell | 1778 | 74 | 24.027027 |
| 2043 | 1960 | Bill | Russell | 1868 | 78 | 23.948718 |

This looks much better. Looking at these numbers, they seem potentially accurate, but those numbers are really high by today's standards, if you are familiar with the current NBA. Did players really get that many rebounds in the '60s? Let's dive in and see if there are trends around rebounding over the years. To do so, the first step is to load a plotting library and start to play around with it.

BASIC PLOTTING WITH SEABORN

First, let's begin with a boxplot of rebounds. There are ways to do this directly in pandas (`nba.boxplot(column=["rebounds"])`), but we are going to use the Seaborn library which is a little more powerful and user friendly. It is using the same matplotlib library that pandas is, but it wraps it in nicer functions. The following shows how to produce a boxplot in Seaborn:

```
sns.boxplot(data=nba.reboundsPerGame)
```



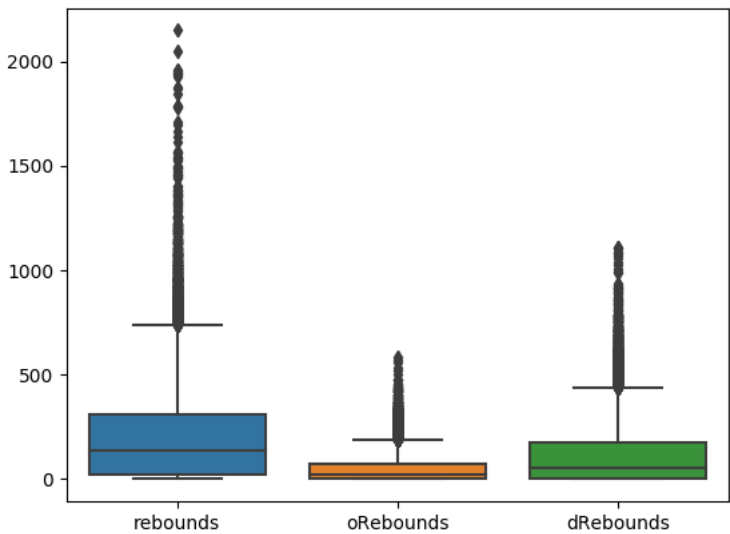
Depending on your environment, running a command like the one above may bring up the plot automatically for viewing. Or it may just prepare it in the background and wait for you to tell it to show it or save it to a file, etc. This can be done with the matplotlib library that we imported earlier as `plt`. With this, we can show the current plot or save it to a file:

```
# Show the current plot
plt.show()

# Save the current plot to a file
plt.savefig("boxplot_reboundsPerGame.png")
```

If we want to do a box plot of multiple columns we can also do that. Here are rebounds, offensive rebounds, and defensive rebounds shown together:

```
sns.boxplot(data=nba[["rebounds", "oRebounds", "dRebounds"]])
```

Obviously, we would want to clean up the presentation of these graphics with better labels, titles, etc. There are many resources on the internet to help understand the parameters that can be passed to these functions. Please go check them out and see what you can do. A major part of learning to use any kind of libraries, and this is especially true with data science libraries, is developing the skill of finding useful information on the internet.

REBOUNDS OVER TIME - APPROACH 1: FACETGRID


With some basic plotting in place, we are ready to revisit the question of whether rebounding trends have changed over time.

One approach could be to use "facets," and put a whole bunch of small boxplots all next to each other. Seaborn has a facet grid function that makes this fairly easy. It might be too much to have a separate plot for every year from 1930-2010, so let's focus just on the '80s first.

```
# Get a subset of the data where the year is between 1980 and 1990
eighties = nba[(nba.year >= 1980) & (nba.year < 1990)]
```


Now let's play around with this to see what we want each facet or mini-plot to look like:

```
sns.boxplot(eighties["reboundsPerGame"], orient="v")
```



That seems to look ok, so now we will set up a FacetGrid and map this function for each facet.

```
grid = sns.FacetGrid(eighties, col="year")
grid.map(sns.boxplot, "reboundsPerGame", orient="v")
```



This is nice and contains a lot of information. But it is really hard to see anything related to the trend we want, namely rebounds from the '60s and '70s versus today. So I don't like it. At this point let's abandon the facet grid approach, and instead, perhaps we could just plot a single point per year, like the median number of rebounds for all players for that year, and look for trends in that.

Instructor Tip:

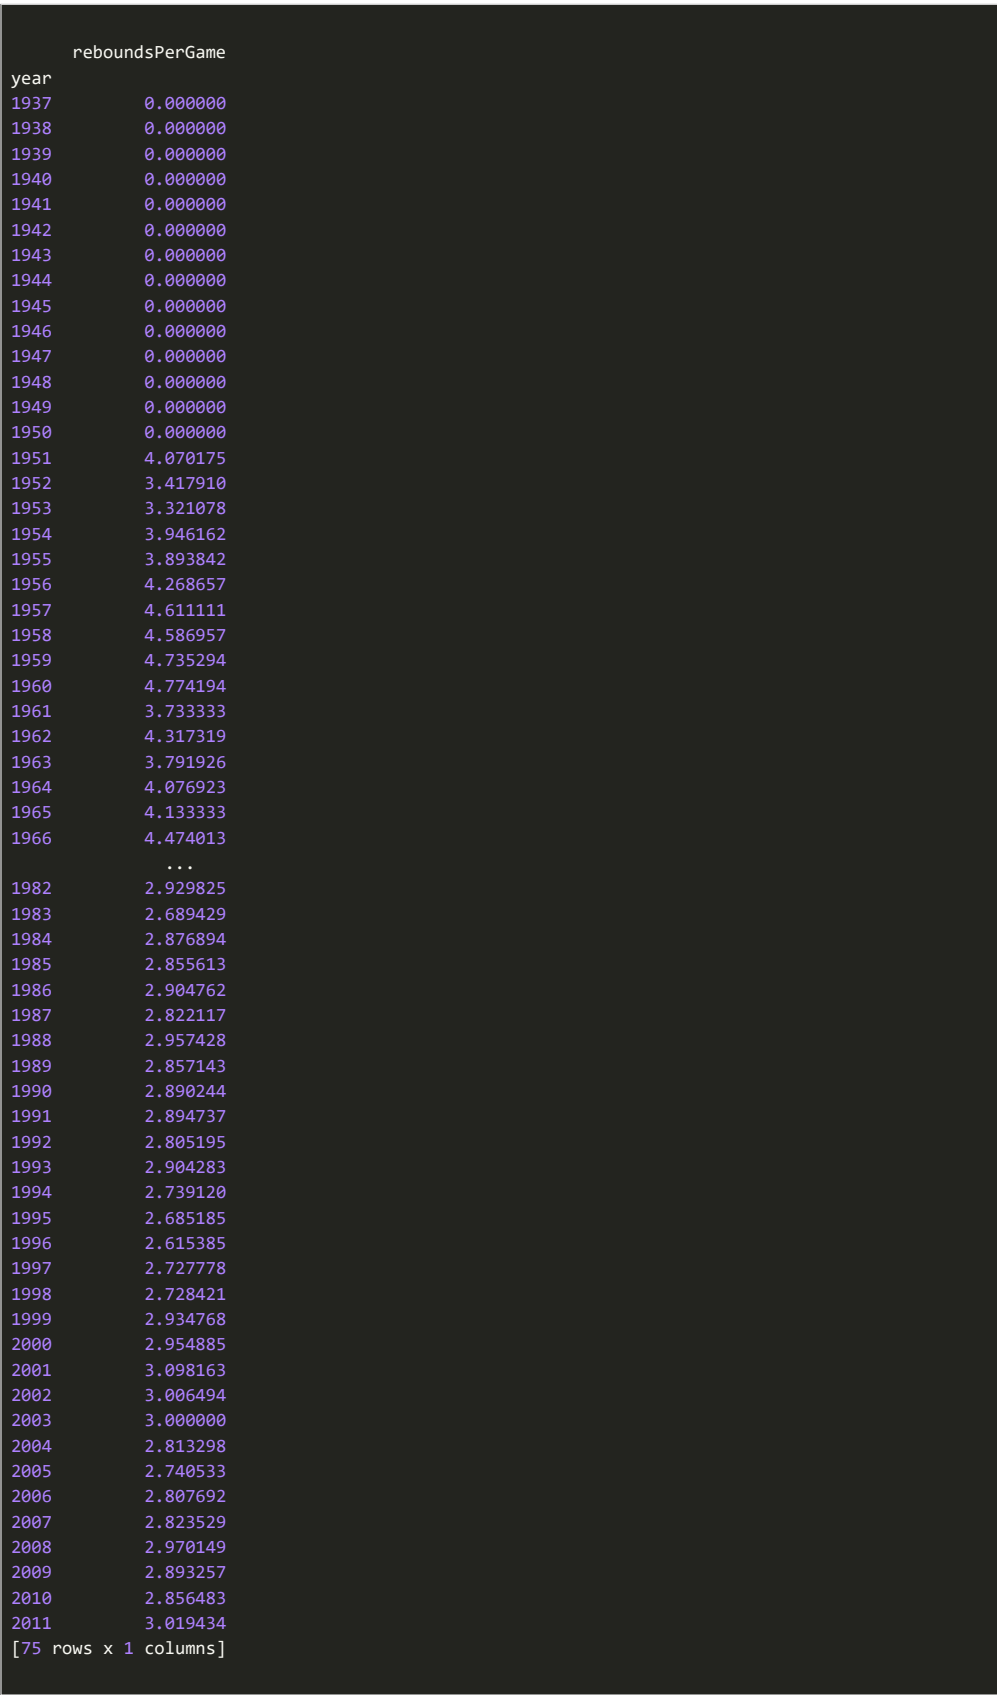
Whoa! Why did we even do these previous steps if we are just going to abandon them?

It turns out that data science is a process of discovery, trial, and error. Sometimes we have an idea, we try it and find out we don't like it, so we have to consider other options. Most tutorials out there (including this one) mostly show the finished product of someone else's discovery process. This is hard, because the important thing for you to learn is actually the process, not the end result.

REBOUNDS OVER TIME - APPROACH 2: GROUPING BY YEAR

To plot rebounds per year, we first need to group the statistics by year. When we do so, we need to specify how we want to aggregate the data of that year. In this case, we'll use the median of the reboundsPerGame.


```
nba_grouped_year = nba[["reboundsPerGame", "year"]].groupby("year").median()
print(nba_grouped_year)
```



Notice that we assigned this to a new variable `nba_grouped_year` so that we can work with this new version of the dataset that is oriented differently.


In order to plot this data, we need to change the index to be the year now, rather than the id that it was previously. Then we can plot it along with a linear regression line as follows:

```
nba_grouped_year = nba_grouped_year.reset_index()
sns.regplot(data=nba_grouped_year, x="year", y="reboundsPerGame")
```

It looks like there are a lot of years where rebounds must not have been tracked (at least in this dataset), so let's remove any years where the median was 0. This time, let's also put a title on the plot.

```
nba_grouped_year = nba_grouped_year[nba_grouped_year["reboundsPerGame"] > 0]
sns.regplot(data=nba_grouped_year, x="year", y="reboundsPerGame").set_title("Median rebounds per
```




Judging from this plot, it looks like there has certainly been a difference in the rebounding from the '60s versus later years.

One question we might ask though is, could this be the result of the fact that we just used the median? Maybe there are a lot more players now, and the top rebounders are still just as productive. We could repeat the previous steps, but this time use the max instead of the median.

```
nba_grouped_year = nba[["reboundsPerGame", "year"]].groupby("year").max()
nba_grouped_year = nba_grouped_year.reset_index()

# Remove the zeros
nba_grouped_year = nba_grouped_year[nba_grouped_year["reboundsPerGame"] > 0]
sns.regplot(data=nba_grouped_year, x="year", y="reboundsPerGame").set_title("Max rebounds per ye
```



SUMMARIZING IN MORE COMPLICATED WAYS

From the previous plot, we can still see a similar trend, which makes us feel a little better about our conclusion. However, this summary is still a little bit troubling, because it could be skewed a lot by the top rebounder of that year. Perhaps that rebounder was a major outlier from the rest of

the league. Another way to consider this is to find the top 10 rebounders of the year and look at their median.

```
# Get the top 10 rebounders per year
nba_topRebounders_perYear = nba[["reboundsPerGame", "year"]].groupby("year")["reboundsPerGame"].nlargest(10)

# Get the median of these 10
nba_topRebounders_median_perYear = nba_topRebounders_perYear.groupby("year").median()


# Put year back in as a column
nba_topRebounders_median_perYear = nba_topRebounders_median_perYear.reset_index()

# Again no zeros...
nba_topRebounders_median_perYear_noZeros = nba_topRebounders_median_perYear[nba_topRebounders_median_perYear["reboundsPerGame"] != 0]

# Now plot
sns.regplot(data=nba_topRebounders_median_perYear_noZeros, x="year", y="reboundsPerGame").set_title("Median of Top 10 Rebounders Each Year")
```

Instructor Tip:

Those are some pretty long variable names! It's up to you what variable names you pick. I often like to choose names that are overly long, so they can be descriptive to a fault. If you are using an editor that autocompletes variable names for you, it's not a problem. But you might choose something a little easier to work with, especially if your editor doesn't autocomplete them.



This feels like a more accurate summary of the top rebounders each season, and it seems to help us answer our original question about rebounding trends. From this graphic we can show that the amount of rebounds per game among the top rebounders fluctuates a little, and peaked around 1970.

V. Concluding Thoughts

This tutorial has walked through the steps of experimenting with rebounds per game, and following questions that might arise in that area.

From this point, we could go on to ask many other questions about this data. What interesting trends can we learn about other statistics? What can we learn about certain players? Are there common characteristics about people from certain eras, hometowns, positions, teams, etc.

Please keep in mind that discovering functions that are available and how to use them is a major discovery process. Following a tutorial like this is really easy, because the work has been done for you of thinking about different options, trying functions that didn't work out well, trying something else, getting errors, looking up solutions on the internet, etc.

This process of discovery, trial, and error takes a lot of work! And this is the real skill that data scientists need. They need to be able to ask questions, like "That's weird, I wonder why ..." and they need to say, "I bet there is a way to put two things on the same plot, I wonder how you do that...". Then you search the internet and figure it out.